


Deep Learning & Transfer Learning

A Comparative Analysis of Image Classification on CIFAR-10


Baseline



Deep Scaling



Transfer Learning



Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck


The Baseline: Simple CNN

Model Stats

Total Parameters: 315,722



Accuracy: ~64% - 68%

Structure: 2 Conv Blocks + Dense





| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------------|---------|
| conv2d_140 (Conv2D) | (None, 30, 30, 32) | 896 |
| max_pooling2d_84 (MaxPooling2D) | (None, 15, 15, 32) | 0 |
| conv2d_141 (Conv2D) | (None, 13, 13, 64) | 18,496 |
| conv2d_141 (Conv2D) | (None, 13, 13, 64) | 18,496 |
| max_pooling2d_85 (MaxPooling2D) | (None, 6, 6, 64) | 0 |
| flatten_38 (Flatten) | (None, 2304) | 0 |
| dense_76 (Dense) | (None, 128) | 295,840 |
| dense_77 (Dense) | (None, 10) | 1,290 |


Total params: 315,722 (1.28 MB)
Trainable params: 315,722 (1.29 MB)
Non-trainable params: 0 (0.00 B)





Escalation: Deep CNN Architecture




Deeper CNN Architecture




deep_baseline(650K): F1=32, F2=64, F3= 128, F4= 256, Dense=256, Batch=32, Epochs= 50 46 LR= 0.001± 0.00003



deep_small(165K): F1=32, F2=64, F3= 128, F4= 256, Dense=256, Batch=32, Epochs=50-35, LR = 0.0005 0.00001




Simple CNN vs Deeper CNN



Transfer Learning: Pretrained Models

EfficientNet B0



EfficientNetB0

Retrain almost 0% of the layers:

| Layer (type) | Output Shape | Param # |
|--|---------------------|-----------|
| input_layer_20 (InputLayer) | (None, 32, 32, 3) | 0 |
| resizing_10 (Resizing) | (None, 240, 240, 3) | 0 |
| lambda_10 (Lambda) | (None, 240, 240, 3) | 0 |
| efficientnetb0 (Functional) | (None, 7, 7, 1280) | 4,049,571 |
| global_average_pooling2d_10 (GlobalAveragePooling2D) | (None, 1280) | 0 |
| dense_20 (Dense) | (None, 128) | 163,968 |
| dense_21 (Dense) | (None, 10) | 1,290 |

Total params: 4,214,829 (16.08 MB)

Trainable params: 165,258 (645.54 KB)

Non-trainable params: 4,049,571 (15.45 MB)

accuracy:
~92 %

Training time:
7min 46s

- Same architecture
- difference in “Trainable params”, training time and accuracy

retrain 20% of the layers:

| Layer (type) | Output Shape | Param # |
|--|---------------------|-----------|
| input_layer_68 (InputLayer) | (None, 32, 32, 3) | 0 |
| resizing_30 (Resizing) | (None, 240, 240, 3) | 0 |
| lambda_29 (Lambda) | (None, 240, 240, 3) | 0 |
| efficientnetb0 (Functional) | (None, 7, 7, 1280) | 4,049,571 |
| global_average_pooling2d_33 (GlobalAveragePooling2D) | (None, 1280) | 0 |
| dense_66 (Dense) | (None, 128) | 163,968 |
| dense_67 (Dense) | (None, 10) | 1,290 |

Total params: 4,214,829 (16.08 MB)

Trainable params: 2,470,554 (9.42 MB)

Non-trainable params: 1,744,275 (6.65 MB)


Training time:
11min 43s

accuracy:
~95 %

EfficientNetB0


retrain 20% of the layers:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| airplane | 0.96 | 0.95 | 0.95 | 1000 |
| automobile | 0.97 | 0.97 | 0.97 | 1000 |
| bird | 0.93 | 0.94 | 0.94 | 1000 |
| cat | 0.88 | 0.91 | 0.89 | 1000 |
| deer | 0.95 | 0.94 | 0.95 | 1000 |
| dog | 0.92 | 0.89 | 0.91 | 1000 |
| frog | 0.96 | 0.98 | 0.97 | 1000 |
| horse | 0.96 | 0.96 | 0.96 | 1000 |
| ship | 0.97 | 0.97 | 0.97 | 1000 |
| truck | 0.96 | 0.96 | 0.96 | 1000 |
| accuracy | | | 0.95 | 10000 |
| macro avg | 0.95 | 0.95 | 0.95 | 10000 |
| weighted avg | 0.95 | 0.95 | 0.95 | 10000 |




EfficientNetB0

retrain 20% of the layers:




Predicting the class of our own pet

Original picture 3605x3605 px:




Resized picture 32x32 px:



1/1 ————— 0s 48ms/step
predicted class: cat

Conclusion: Performance vs. Cost



Final Verdict:

Transfer Learning achieves a decisive victory in accuracy (95%), offering the best performance-to-effort ratio for general classification.

However, this comes with a 7M parameter footprint and increased pre-processing overhead.